

Veon Almeida

Chicago, IL | veonalmeida2003@gmail.com | (312) 394-9398 | linkedin.com/in/veon-almeida | github.com/Aries2003

PROFESSIONAL SUMMARY

AI Master's candidate who builds software and GenAI agents that automate manual processes, owning each feature from start to finish: learning the business context, shipping a working version, and rigorously self-testing before handoff. Uses AI coding agents as force multipliers while reviewing every line, writes clear documentation, and communicates progress and changing estimates early.

EDUCATION

DePaul University

M.S. in Artificial Intelligence — GPA: 3.75/4.0

Relevant Coursework: Machine Learning I (DSC 445), Deep Learning (CSC 383)

Chicago, IL
Expected June 2027

TECHNICAL SKILLS

AI-Augmented Development: Claude Code, Cursor, GitHub Copilot [CONFIRM], Prompt Engineering, Reviewing & Testing AI-Generated Code

LLMs & AI: Transformer Architecture, Embeddings, Tokenization, RAG, LangChain, LangGraph, Chroma Vector DB, LLM Evals, Hugging Face, PyTorch

Languages & Web: Python, TypeScript/JavaScript, React, HTML, CSS, SQL, C++, C#, REST APIs

Engineering Workflow: Terminal/CLI, Git (command line), GitHub Actions CI/CD, Docker, Linux, Bash, AWS, GCP

EXPERIENCE

RAISE Lab (Responsible AI Systems and Societal Experiences), DePaul University

Research Assistant

Chicago, IL
June 2026 – Present

- Sat with researchers to understand how they actually labeled speech data, then built an automated labeling pipeline with Hugging Face Transformers and rule-based validation, cutting manual effort by 60% and speeding up curation by 30%.
- Built Python evaluation pipelines on Linux benchmarking 8+ deep learning models across 100+ hours of data, used by the lab across 3 research projects.
- Engineered GAN-based synthetic data generation in PyTorch to fill gaps in scarce training data, improving downstream model robustness by 12%.
- Wrote readable, reusable tooling with clear documentation and demoed results to faculty, reducing new-tool onboarding time by 20%.

Tera Bhai Entertainment

AI Developer

Remote

Oct 2022 – May 2024

- Built an NLP-driven dialogue and AI behavior system in Unity (C#) across 50+ interaction scenarios, working directly with designers and iterating on features from real playtest feedback.
- Picked up Unreal Engine 4 and Godot from documentation to ship AI features across all three engines.

TDS Corp

Machine Learning Engineer

Remote

Mar 2021 – Oct 2021

- Built a modular API testing automation framework that reduced regression bugs by 60%, making QA a built-in step of every monthly release.
- Streamlined ETL preprocessing with pandas and NumPy, cutting pipeline runtime by 35%, and built SQL dashboards giving stakeholders visibility into model metrics.

PROJECTS

GenAI Agent: LangGraph Ordering Assistant (Used by 120+ Users) | Tech Stack: Python, LangGraph, LangChain, OpenAI API, React, Docker, GitHub Actions

- Owned the project end to end: scoped the conversation flow, built a multi-step agentic workflow with validation checkpoints, and shipped a live web demo with 99% uptime across 120+ user sessions.
- Instrumented session analytics to see how people actually used it, turning observed behavior into 3 iterative releases, and added role-based access controls for secure testing.

GenAI Document Q&A Agent (RAG) | Tech Stack: Python, LangChain, Chroma Vector DB, Hugging Face Embeddings, LLM Evals

- Built a retrieval-augmented generation assistant with sub-second retrieval over a 50,000-document corpus and source-grounded answers.
- Designed a 200-question evaluation set to catch where the LLM quietly produced wrong answers, iterating on chunking and prompts to reach a 93% relevance rate.

Rep.IO – Real-Time Exercise Tracking | Tech Stack: Python, MediaPipe, OpenCV, Pose Estimation

- Built a real-time computer vision app that tracks body keypoints from live video at 30 FPS to count reps and give form feedback.
- Optimized the frame-processing loop so feedback stays real-time on consumer hardware, with joint-angle analysis flagging incorrect form.

County-Level Election Prediction | Tech Stack: Python, pandas, XGBoost, Optuna, SHAP

- Built an end-to-end data and modeling pipeline across 3,000+ counties, reaching 90%+ accuracy and explaining predictions to non-technical audiences with SHAP.
- Tuned XGBoost with Optuna and compared it against MLP and stacked ensemble models to pick the most reliable approach.